

TITLE OF PRESENTATION

THIS IS A MORE DETAILED SUBTITLE

ROBERT MOSS

STANFORD UNIVERSITY

mossr@cs.stanford.edu

TITLE OF SLIDE

Example equation and algorithm block with `julia` verbatim in a slide.¹

$$\underbrace{V(s)}_{\text{state value}} \leftarrow \max_a \overbrace{\left(\underbrace{R(s, a)}_{\text{reward}} + \underbrace{\gamma V(s')}_{\substack{\text{discounted value} \\ \text{of next state}}} \right)}^{\text{estimate of optimal discounted value}}$$

```
# Julia function
V(s) = max(a->R(s,a) + γ*V(s'), A)
```

¹Mykel J. Kochenderfer and Tim A. Wheeler. *Algorithms for Optimization*. MIT Press, 2019.

JULIA CONSOLE

Example Julia code executing at LaTeX compilation time.

```
julia> using LinearAlgebra
```

```
julia> A = Matrix{Int}(I, 3, 3)
```

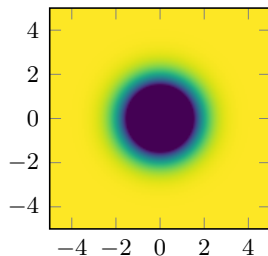
```
3×3 Array{Int64,2}:
```

```
 1  0  0  
 0  1  0  
 0  0  1
```

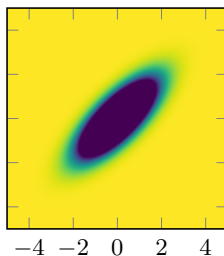
PLOTTING

Plot using `PGFPlots.jl`² directly in the TeX file.

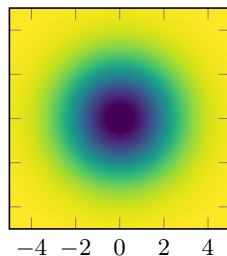
$$\mu = [0, 0]$$
$$\Sigma = \begin{bmatrix} 1 & 0 \\ 0 & 1 \end{bmatrix}$$



$$\mu = [0, 0]$$
$$\Sigma = \begin{bmatrix} 1 & 0.75 \\ 0.75 & 1 \end{bmatrix}$$



$$\mu = [0, 0]$$
$$\Sigma = \begin{bmatrix} 3 & 0 \\ 0 & 3 \end{bmatrix}$$



²<https://github.com/JuliaTeX/PGFPlots.jl>

TIKZ FIGURES

For a neural network with one hidden layer, the intermediate hidden units are $h_j = \sigma(\mathbf{v}_j \cdot \phi(x))$ where $\sigma(z) = (1 + e^{-z})^{-1}$, producing output score = $\mathbf{w} \cdot \mathbf{h}$:

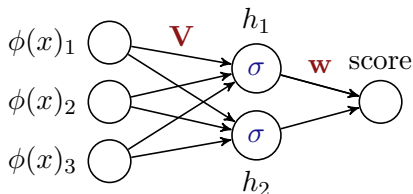


Figure: A one-layer neural network with $\phi(x) \in \mathbb{R}^3$ and $\mathbf{h} \in \mathbb{R}^2$ using the logistic activation function σ .

```
σ(z) = 1/(1 + exp(-z)) # sigmoid

function neural_network(x, V, w, φ, g::Function=σ)
    h = map(v_j -> g(v_j · φ(x)), V)
    w · h
end
```

ALGORITHMS

- Monte Carlo tree search (MCTS) is an anytime algorithm that uses rollouts of a random policy to estimate the value of each state-action node in the tree.³
- There are four main stages in each simulation: *selection*, *expansion*, *rollout* (or *simulation*), and *backpropagation*.
- The tree \mathcal{T} is iteratively expanded and the policy improves over time as the algorithm balances exploration with exploitation of the state and action spaces.

Algorithm 1 Top-level Monte Carlo tree search algorithm.

```
function MONTECARLOTREESearch( $s, d$ )
  loop
    SIMULATE( $s, d$ )
  return  $\arg \max_{\bar{a} \in A(s)} Q(s, \bar{a})$ 
```

Algorithm 2 Monte Carlo tree search simulation.

```
function SIMULATE( $s, d$ )
  if  $d = 0$ 
    return 0
  if  $s \notin \mathcal{T}$ 
     $\mathcal{T} \leftarrow \mathcal{T} \cup \{s\}$ 
     $N(s) \leftarrow N_0(s)$ 
    return ROLLOUT( $s, d$ )
   $N(s) \leftarrow N(s) + 1$ 
   $\bar{a} \leftarrow \text{SELECTACTION}(s)$  ▷ selection
   $(s', r) \leftarrow \text{DETERMINISTICSTEP}(s, \bar{a})$  ▷ expansion
   $q \leftarrow r + \gamma \text{SIMULATE}(s', d - 1)$  ▷ simulation/rollout
   $N(s, \bar{a}) \leftarrow N(s, \bar{a}) + 1$ 
   $Q(s, \bar{a}) \leftarrow Q(s, \bar{a}) + \frac{q - Q(s, \bar{a})}{N(s, \bar{a})}$  ▷ backpropagation
  return  $q$ 
```

³Rémi Coulom. “Efficient selectivity and backup operators in Monte-Carlo tree search”. In: *International Conference on Computers and Games*. Springer. 2006, pp. 72–83.

DEFINITION BLOCKS

Example use of a `definitionblock` environment with accompanying Julia code.

Definition: score. The score on an example (x, y) is $\mathbf{w} \cdot \phi(x)$, how *confident* we are in predicting $+1$. Score is a weighted combination of features:

$$\mathbf{w} \cdot \phi(x) = \sum_{j=1}^d w_j \phi(x)_j$$

```
score(x, w, φ) = w·φ(x)
```

TABLES

Example table with footnotes.

Table: Algorithm Hyperparameters

Hyperparameter	Value
episodes [*]	5000
maximum tree depth d_{\max} (i.e. number of waypoints) [*]	12
rollout depth d [†]	12
exploration constant c	10
progressive widening k	10
progressive widening α	0.3

^{*} Used by all algorithms.

[†] Used by MCTS and direct Monte Carlo.

EXAMPLE: BULLET POINTS

- Bullet point
 - Sub-bullet point
- Another bullet point
 - Another sub-bullet point
 - ▶ Another level

REFERENCES

- Coulom, Rémi. “Efficient selectivity and backup operators in Monte-Carlo tree search”. In: *International Conference on Computers and Games*. Springer. 2006, pp. 72–83.
- Kochenderfer, Mykel J. and Tim A. Wheeler. *Algorithms for Optimization*. MIT Press, 2019.